


a)



6

All u-threads are multiplexed on one OS process.

WAIT(s): add calling thread to s-queue on need to block then call BLOCK(u-thread) + schedule

SIGNAL(s): if s has a queue, remove one thread and call UNBLOCK(u-thread).

b) i) thread management is in the OS.

the OS schedules threads.

the RTS is involved only in create/kill (thread) - making threads known to the OS.

eg k-thread-id := create(thread) system call

a blocking call to WAIT(s): manage s-queue as before and call OS block(k-thread-id).

an unblocking call to SIGNAL(s): manage s-queue as before and call OS unblock(k-thread-id).

5

Assumption - no limit to # of threads - can use one k-thread for each u-thread.

(ii) OS guarantees that some fixed number of k-threads per address space can RUN.

Assumption: OS tells RTS if a thread makes a blocking system call & passes its context. The RTS passes the context of the runnable thread to be run in its place.

Blocking at user level on WAIT(s) requires s-data value & queue management & OS calls to block this thread & set up another. An unblocking call to

5

SIGNAL(s): manage s-queue & change state of unblocked thread at user level

4

c) Semaphore data structures are shared & read/write not atomic. WAIT & SIGNAL must be made atomic - one (monitor) lock/module or a lock (boolean) per semaphore. Threads spin (for a short time) on these locks on which they are waiting to block

(20)