

Buffer only - write through synchronously to disc
- read: block may be in memory from a previous read, if not, read disc.

2 Disc block cache: read as ⁹, write to block in memory
write to disc when ~~block~~^{buffer} needed to store some other block
(and probably periodically eg UNIX)

Pro - fast response to write

Con - "done" to application on write does not \Rightarrow data on disc
 \therefore can't use for Tx. system. May be lost on a crash.
Writes could be re-ordered + interaction with crashes.

Concurrency Control

App. (hyp-damn) access to buffer must wait for device driver to fill buffer before a read.

Access to all data such as free list + block headers must take place under mutual exclusion to avoid race condition / dual allocation etc.

6 Must wait for a dirty block to be written out before filling

typical header of a disc block buffer:

filesystem/block (VID)

owner: PID.

dirty (changed)

valid? (being filled)
_{emptied}

locked (for investigation)

links for hash chain + free list

eg: Read (bytes) \rightarrow block ID \rightarrow in cache? \rightarrow YES \rightarrow valid \rightarrow read
 \rightarrow N. \rightarrow wait

(waiting involves sleep + wakeup when event signalled)

N: look free chain - try first block - dirty? \rightarrow Y. write out try next.
 \rightarrow N allocate

\rightarrow read (mark invalid)
wait (valid)

Write (bytes) \rightarrow block ID \rightarrow in cache? \rightarrow YES \rightarrow valid \rightarrow write
 \rightarrow N. \rightarrow write mark dirty
N) wait \rightarrow

9 N: get a new buffer as for read.

< expand >.