

strong - invokers always see the same version of the replicas  
 weak - invokers may see different versions but must have eventual convergence.

First consider one request to a replica manager.

Use quorum assembly (of a read or write quorum) where  
 $WQ > \frac{n}{2}$ ,  $RQ + WQ > n$ .

Having assembled a quorum, update all replicas to the latest version then read or write. Propagate writes in background.

Failures: run two or three phase commit to achieve quorum update. If request is abort due to a failed member, attempt to get another replica in quorum and try 2PC again. If commit manager fails, group may attempt to recover (as in 3PC) or invoker may detect & retry some other replica manager.

Concurrency: suppose 2 invocations are in progress concurrently. One may get a quorum and delay other. BUT both could each lock half & deadlock - must allow for this - e.g. the assemblers may detect & latest timestamp back off.....

[Solutions will take longer to say all this - outline given].

Structured group, single coordinator

Not appropriate for widely distributed replicas - better to invoke a nearby one. Single coordinator is a single point of failure (leading to need for election of a new coordinator) and a potential bottleneck - but this orders requests so avoids concurrent updates - but we can handle concurrent update requests as above.

Strong consistency not appropriate for apps. generating large numbers of concurrent invocations & requiring timely service.