

Solution notes

Data Structures and Algorithms 2005 (MR) Paper 6 Question 1, Paper 13 Question 1

This question draws on basic knowledge of many algorithms covered in the course, together with an ability to reason about algorithm costs.

- (a) Vector `count[1] ... count[N]` holding the counts. A vector `index[1] ... index[N]` holding counter numbers in increasing order of value. So `count[index[1]]` is the smallest counter value (`=mincount`) and `count[index[N]]` is the largest counter value (`=maxcount`). A vector `inverse` which is the inverse of `index`, ie `index[inverse[i]]=i` so that when `increment(i)` is called it can find which element of `index` is affected. When `increment(i)` is called, `count[i]` is incremented and then the elements of `index` and `inverse` possibly changed to maintain the ordering. Some adjacent elements of `index` may have to be swapped, but not many if `N` is known to be about 10. Alternatively, do not keep the counters in sorted order but look through all `N` counters when one with the minimal value is incremented to see if `mincount` has to change. If a counter with value `maxcount` is incremented then so is `maxcount`.
- (b) If `N` is about 10^6 the pairwise exchanges may become too expensive. For instance incrementing counter 1 when all counter are zero would cause 10^6-1 exchanges. This could be much improved by hold the counts in a priority queue based on a binary heap as in `heapsort`. `count[index[1]]` would still be `mincount`, but otherwise the counters would not be fully sorted. The number of swaps needed when a counter is incremented would now be no more than $\log_2(N)$ (~ 20) but typically much less. If the incremented counter is larger than `maxcount`, `maxcount` would be updated. If `count[index[1]]` changes, `mincount` would be updated. Alternatively, if `maxcount-mincount` is known not to be too large, the following scheme could be used. Allocate a vector `v` of integers somewhat larger than the expected maximum difference of `maxcount` and `mincount`. This vector will be used as a circular buffer containing the count of how many counters have each possible value. `v[mincount mod size]` will be the number of counters with value `mincount`, `v[maxcount mod size]` will be the number of counters holding value `maxcount`, etc. In general, `v[count[i] mod size]` will hold the number of counters with value `count[i]`. When counter `i` is incremented `v[count[i] mod size]` is decremented and `v[(count[i]+1) mod size]` is incremented. With this scheme it is easy to determine when to increment `mincount`. Dealing with `maxcount` is already easy.
- (c) At the end of the run the average value of each counter will be about 1000. So halfway through the run they will have values around 500. In a sorted list of these count values we can expect long sequences of equal values, probably in the region of 2000 to 10000 in length. We can thus expect the number of swaps per counter increment

when using algorithm (a) to between 1000 and 5000, compared with probable 2 or 3 swaps needed for algorithm (b) The improvement therefore somewhere around $1000/3$ to $5000/2$, or 300 to 2500. Certainly significant. Better students may even produce more accurate estimates, but probably not many.